

DSC680_Project1_jMadsen

April 9, 2023

```
[1]: # Justin Madsen

# import the boys
from datetime import datetime, timedelta
from prophet import Prophet
from prophet.plot import plot_plotly, plot_components_plotly, \
    plot_cross_validation_metric
from prophet.diagnostics import cross_validation, performance_metrics
import matplotlib.pyplot as plt
import numpy as np
import pandas as pd
import pandas_datareader.data as pdr
import plotly.graph_objects as go
import yfinance as yf
```

```
[2]: # yahoofinance changes made datareader brick, so the yahoofinance library \
    included a function to bypass
# the changes to the website.
yf.pdr_override()

# create a variable that hold's today's date, then convert to a d/t group
today = datetime.today().strftime('%Y-%m-%d')

# here we use datareader to ingest the historic ethereum prices
eth_ingest = pdr.get_data_yahoo(['ETH-USD'], start=datetime(2016, 1, 1), \
    end=today)
```

[*****100%*****] 1 of 1 completed

```
[3]: # did it work?
eth_ingest.head(2)
```

```
[3]:
```

	Open	High	Low	Close	Adj Close	\
Date						
2017-11-09	308.644989	329.451996	307.056000	320.884003	320.884003	
2017-11-10	320.670990	324.717987	294.541992	299.252991	299.252991	

Volume

```
Date
2017-11-09  893249984
2017-11-10  885985984
```

```
[4]: # let's save the data as a csv incase the yfinance page changes and we can't
      ↪pull the data
      eth_ingest.to_csv('eth.csv')
```

```
[5]: # reingest the data
      eth_df = pd.read_csv('eth.csv')
```

```
[6]: # did it work?
      eth_df.head(2)
```

```
[6]:
```

	Date	Open	High	Low	Close	Adj Close	\
0	2017-11-09	308.644989	329.451996	307.056000	320.884003	320.884003	
1	2017-11-10	320.670990	324.717987	294.541992	299.252991	299.252991	


```

      Volume
0  893249984
1  885985984
```

```
[7]: # how much data are we looking at?
      eth_df.shape
```

```
[7]: (1977, 7)
```


```
[8]: # let's take a look at what the data looks like
      eth_df.describe().T
```

```
[8]:
```

	count	mean	std	min	25%	\
Open	1977.0	1.153107e+03	1.163814e+03	8.427969e+01	2.173270e+02	
High	1977.0	1.189362e+03	1.199071e+03	8.534274e+01	2.221826e+02	
Low	1977.0	1.112480e+03	1.123251e+03	8.282989e+01	2.096381e+02	
Close	1977.0	1.153654e+03	1.163301e+03	8.430830e+01	2.171830e+02	
Adj Close	1977.0	1.153654e+03	1.163301e+03	8.430830e+01	2.171830e+02	
Volume	1977.0	1.267252e+10	1.058283e+10	6.217330e+08	4.709988e+09	

	50%	75%	max
Open	5.893787e+02	1.746926e+03	4.810071e+03
High	6.085830e+02	1.806539e+03	4.891705e+03
Low	5.685964e+02	1.691658e+03	4.718039e+03
Close	5.896632e+02	1.752045e+03	4.812087e+03
Adj Close	5.896632e+02	1.752045e+03	4.812087e+03
Volume	1.029222e+10	1.774097e+10	8.448291e+10

```
[9]: # let's plot all the data to look at distributions
eth_df.hist(bins=50, figsize=(20,15))
plt.savefig('eth_hist.png')
plt.show()
```



Close and Adj close look pretty similar. Let's do some comparison.

```
[10]: # np.where is useful to create these comparisons in a new column
eth_df['diff'] = np.where(eth_df['Close'] != eth_df['Adj Close'], 'Yes', "No")

[11]: # did it work?
eth_df.head(2)
```

```
[11]:
```

	Date	Open	High	Low	Close	Adj Close	\
0	2017-11-09	308.644989	329.451996	307.056000	320.884003	320.884003	
1	2017-11-10	320.670990	324.717987	294.541992	299.252991	299.252991	

	Volume	diff
0	893249984	No
1	885985984	No

Let's check to see if the Close and Adj Close ever differ

```
[12]: # value_counts() counts each different type of entry in whatever column we
      ↪ specify
      eth_df['diff'].value_counts()
```

```
[12]: No      1977
      Name: diff, dtype: int64
```

Since Close and Adj Close are the same, we can go ahead and drop Adj Close. This column adds no value and just takes up space.

```
[13]: # we can drop the diff column since we already did the comparison
      eth_df = eth_df.drop(['Adj Close', 'diff'], axis = 1)
      eth_df.head(2)
```

```
[13]:      Date      Open      High      Low      Close      Volume
0  2017-11-09  308.644989  329.451996  307.056000  320.884003  893249984
1  2017-11-10  320.670990  324.717987  294.541992  299.252991  885985984
```

```
[14]: # let's check the types of the objects
      eth_df.dtypes
```

```
[14]: Date      object
      Open      float64
      High      float64
      Low       float64
      Close     float64
      Volume     int64
      dtype: object
```

Date is an object. Let's go ahead and convert that to an actual datetime group

```
[15]: eth_df['Date'] = pd.to_datetime(eth_df['Date'], format='%Y-%m-%d')
      eth_df.dtypes
```

```
[15]: Date      datetime64[ns]
      Open      float64
      High      float64
      Low       float64
      Close     float64
      Volume     int64
      dtype: object
```

```
[16]: # numeric_only is used in order to test for correlation with dates since it's
      ↪ not a numerical value
      corr_matrix = eth_df.corr(numeric_only=False)
      corr_matrix
```

```
[16]:
```

	Date	Open	High	Low	Close	Volume
Date	1.000000	0.629282	0.626550	0.633534	0.629352	0.503962
Open	0.629282	1.000000	0.999189	0.998236	0.997744	0.528284
High	0.626550	0.999189	1.000000	0.998055	0.998857	0.538204
Low	0.633534	0.998236	0.998055	1.000000	0.998891	0.509392
Close	0.629352	0.997744	0.998857	0.998891	1.000000	0.525440
Volume	0.503962	0.528284	0.538204	0.509392	0.525440	1.000000

```
[17]: # open, high, low, close are all almost a flat 1.0 in correlation together.
# building a model on those would be ill advised. Let's do date and close.
input_df = eth_df[['Date', 'Close']]
```

```
[18]: # fbprophet requires columns to be named ds and y, so let's rename them in a
      ↪ new df
input_df = input_df.rename(columns = {'Date': 'ds', 'Close': 'y'})
```

```
[19]: # did it work?
input_df.head(2)
```

```
[19]:
```

	ds	y
0	2017-11-09	320.884003
1	2017-11-10	299.252991

```
[20]: # here we'll set an x and y based off the column data
x = input_df['ds']
y = input_df['y']

# using plotly, let's create the figure
fig = go.Figure()

# add the line for history of the price
fig.add_trace(go.Scatter(x=x, y=y))

# let's add a title
fig.update_layout(title_text="Ethereum Price History")
```

```
[21]: # here we'll save the image
fig.write_image('eth_price.png')
```

```
[22]: # create the model using Prophet, and use multiplicative as the seasonality
      ↪ since this is a time series model
forecast_model = Prophet(seasonality_mode="multiplicative")
forecast_model.fit(input_df)
```

```
16:07:10 - cmdstanpy - INFO - Chain [1] start processing
16:07:10 - cmdstanpy - INFO - Chain [1] done processing
```

```
[22]: <prophet.forecaster.Prophet at 0x1c88b9313c0>
```

```
[23]: # let's set a forecast for 1 year
future_dates = forecast_model.make_future_dataframe(periods = 365)
future_dates.tail()
```

```
[23]:          ds
2337 2024-04-03
2338 2024-04-04
2339 2024-04-05
2340 2024-04-06
2341 2024-04-07
```

```
[24]: # Now let's create the forecast off those dates
forecast = forecast_model.predict(future_dates)

# let's grab the forecast dates and their predictions.
forecast[['ds', 'yhat', 'yhat_lower', 'yhat_upper']].tail()
```

```
[24]:          ds          yhat  yhat_lower  yhat_upper
2337 2024-04-03  27.729676 -2252.649441  2156.760446
2338 2024-04-04  23.604549 -2326.711371  2164.035329
2339 2024-04-05  19.530385 -2257.031090  2123.253921
2340 2024-04-06  15.624360 -2250.241658  2106.686575
2341 2024-04-07  11.718896 -2275.002692  2178.200608
```

```
[25]: # using datetime, let's create an object that uses the next day
next_day = (datetime.today() + timedelta(days=1)).strftime('%Y-%m-%d')

# what's the forecast for tomorrow?
forecast[forecast['ds'] == next_day]['yhat'].item()
```

```
[25]: 1388.6476583136105
```

```
[26]: # let's plot the forecast
fig2 = plot_plotly(forecast_model, forecast)
fig2
```

```
[27]: # save the image
fig.write_image('eth_forecast.png')
```

```
[28]: # now let's plot the trend line, as well as the deviation by year and week
plot_components_plotly(forecast_model, forecast)
```

```
[29]: # let's get some cross validation
forecast_model_cv = cross_validation(forecast_model, initial="730 days",
    ↪period='180 days', horizon='730 days')
forecast_model_cv
```

```
0%|          | 0/3 [00:00<?, ?it/s]
```

```

16:07:11 - cmdstanpy - INFO - Chain [1] start processing
16:07:11 - cmdstanpy - INFO - Chain [1] done processing
16:07:11 - cmdstanpy - INFO - Chain [1] start processing
16:07:11 - cmdstanpy - INFO - Chain [1] done processing
16:07:11 - cmdstanpy - INFO - Chain [1] start processing
16:07:12 - cmdstanpy - INFO - Chain [1] done processing

```

```

[29]:
      ds      yhat  yhat_lower  yhat_upper      y      cutoff
0  2020-04-14  118.989870    54.240048    184.604011  157.596390  2020-04-13
1  2020-04-15  120.661759    58.886126    186.419713  153.286896  2020-04-13
2  2020-04-16  122.608373    58.868049    183.331217  172.157379  2020-04-13
3  2020-04-17  126.267111    69.642778    188.718630  171.638580  2020-04-13
4  2020-04-18  131.997180    69.274435    195.063494  186.914001  2020-04-13
...
2185 2023-04-04  6600.230564  4463.294266  8748.476759  1871.005127  2021-04-08
2186 2023-04-05  6628.469267  4535.749458  8778.256973  1909.114014  2021-04-08
2187 2023-04-06  6630.266560  4521.616140  8813.544810  1872.922607  2021-04-08
2188 2023-04-07  6735.758681  4525.768108  8903.723588  1865.636108  2021-04-08
2189 2023-04-08  6860.667871  4649.362796  9109.353520  1849.498169  2021-04-08

```

[2190 rows x 6 columns]

```

[30]: # assign the performance to an object
forecast_model_performance = performance_metrics(forecast_model_cv)
forecast_model_performance

```

```

[30]:
      horizon      mse      rmse      mae      mape      mdape \
0    73 days  1.001598e+05  316.480292  223.292154  0.256312  0.236747
1    74 days  1.054573e+05  324.741970  228.719168  0.260638  0.238435
2    75 days  1.107734e+05  332.826450  234.132244  0.265034  0.240154
3    76 days  1.149114e+05  338.985863  238.993898  0.268755  0.243701
4    77 days  1.188438e+05  344.737241  243.923212  0.272543  0.247385
..
653 726 days  1.078257e+07  3283.682126  2727.939734  1.392605  1.038883
654 727 days  1.081923e+07  3289.260225  2733.217138  1.391730  1.037952
655 728 days  1.084691e+07  3293.464426  2735.884740  1.390520  1.037937
656 729 days  1.087631e+07  3297.924856  2738.301868  1.388662  1.037323
657 730 days  1.091672e+07  3304.046622  2741.854614  1.387753  1.036980

      smape  coverage
0    0.302470  0.342466
1    0.306826  0.328767
2    0.311235  0.315068
3    0.315055  0.301370
4    0.319008  0.287671
..
653  1.149865  0.333333
654  1.149765  0.333333


```

```
655 1.148956 0.333333
656 1.147820 0.333333
657 1.146761 0.333333
```

```
[658 rows x 8 columns]
```

Here we can see the errors greatly increase the further down the forecast window we go. This is expected as it isn't pulling in new data, and only has a few years of data to go off of.

```
[31]: # plot the cv values
fig = plot_cross_validation_metric(forecast_model_cv, metric = 'mape')
```



1 references

How to install fbprophet on Windows - <https://stackoverflow.com/a/64878241>

Facebook.Prophet documentation - <https://facebook.github.io/prophet/docs/diagnostics.html>

<https://github.com/facebook/prophet/blob/main/python/prophet/forecaster.py>

```
[ ]:
```